

MACHINE LEARNING 1
ADVANCED CRIME ANALYSIS
UCL

BENNETT KLEINBERG

18 FEB 2019

MACHINE LEARNING 1

TODAY

- Recap week 1-5
- Intro to machine learning
 - Types of ML
 - Supervised machine learning
 - Step-by-step example
 - Important algorithms

RECAP WEEK 2

APIS

RECAP WEEK 3

WEBSCRAPING

RECAP WEEK 4

TEXT MINING 1

RECAP WEEK 5

TEXT MINING 2

MACHINE LEARNING?

- core idea: a system learns from experience
- no precise instructions

Applications?

WHY DO WE WANT THIS?

Step back...

How did you perform regression analysis in PSM2?

OKAY ...

- you've got one outcome variable (e.g. number of shooting victims)
- and two predictors (e.g. gender of shooter, age)
- typical approach *victims* $gender + age$
- regression equation with intercept, beta coefficients and inferred error term

BUT!

Often we have no idea about the relationships.

- too many predictors
- too diverse a problem
- simply unknown

ML IN GENERAL

- concerned with patterns in data
- learning from data
- more experience results typically in better models
- data, data, data

TYPES OF MACHINE LEARNING

BROAD CATEGORIES

- Supervised learning (today)
- Unsupervised learning (next week)
- Hybrid models
- Deep learning
- Reinforcement learning

DEEP LEARNING

Inspired by the human brain.

- MIT's course website <https://deeplearning.mit.edu/>
- Lex Fridman's courses from MIT → [YouTube](#)

REINFORCEMENT LEARNING

- Excellent YouTube examples from [code bullet](#)
- e.g. [AI Learns to play the Worlds Hardest Game](#)

[Demo](#)

SUPERVISED LEARNING

WTF IS SUPERVISED?

- supervised = labeled data
- i.e. you know the outcome
- flipped logic

Contrary: unsupervised.

CLASSES OF SUPERVISED LEARNING

- classification (e.g. death/alive, fake/real)
- regression (e.g. income, number of deaths)


MINI EXAMPLE

Supervised classification


SIMPLE EXAMPLE

- gender prediction
- based on salary

gend	er sal	ary
1	male	39169
2	male	33620
3	male	33225
4	male	35437
11	female	15039
12	female	13861
13	female	24443
14	female	36744




How to best separate the data into two groups?



CORE IDEA

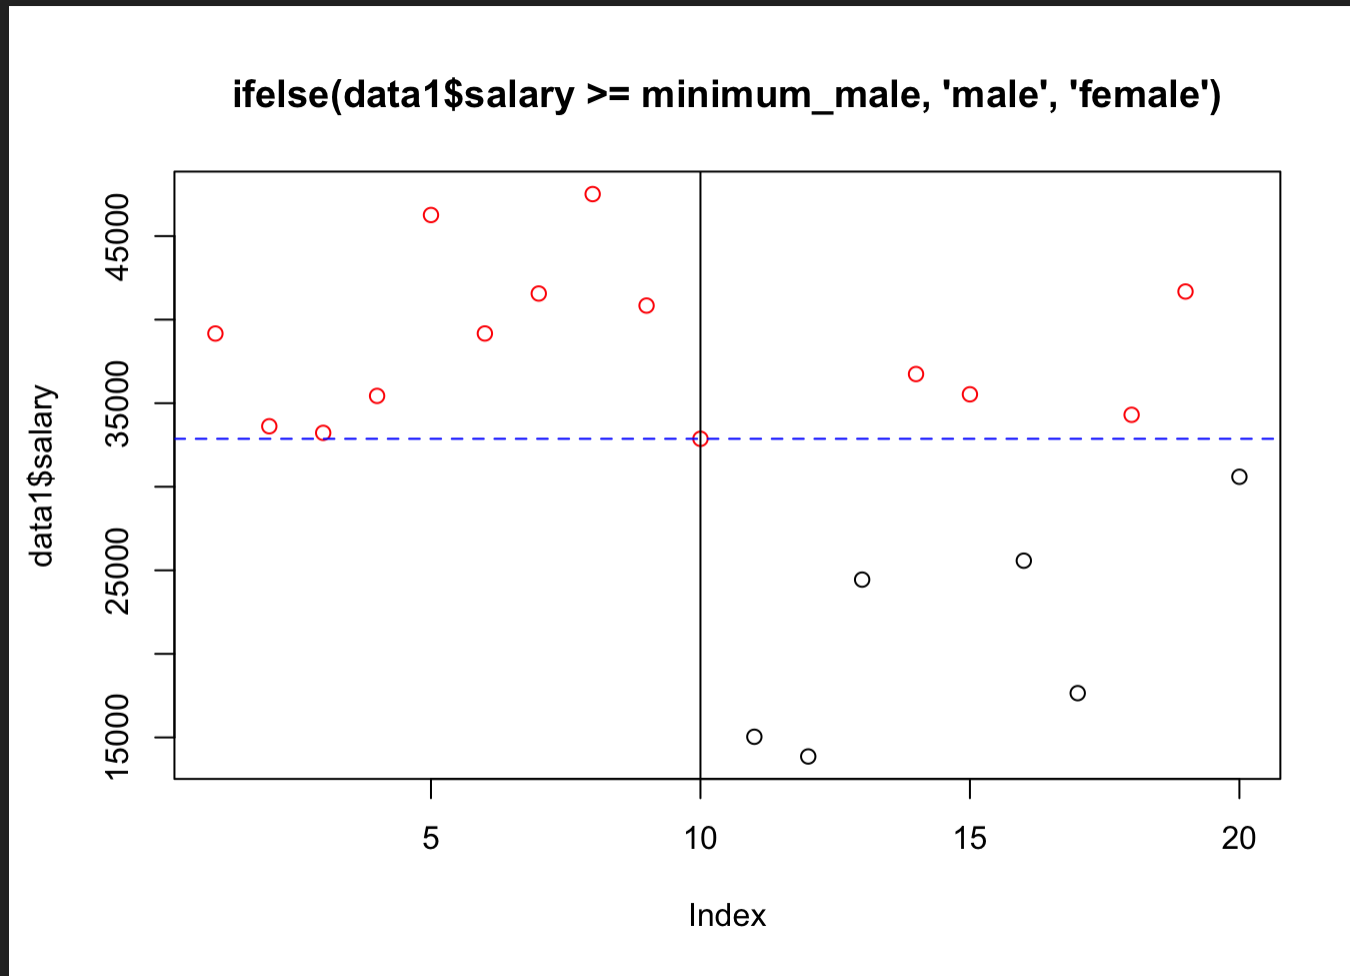
- learn relationship between
 - outcome (target) variable
 - features (predictors)
- “learning” is done through an algorithm
 - simplest algorithm: `if A then B`

IDEA 1: MALE SALARY THRESHOLD




IDEA 1: MALE SALARY THRESHOLD

```
minimum_male = min(data1$salary[data1$gender == 'male']) #32869  
data1$my_prediction = ifelse(data1$salary >= minimum_male, 'male', 'f
```




IDEA 2: FEMALE SALARY THRESHOLD



IDEA 2: FEMALE SALARY THRESHOLD

```
maximum_female = max(data1$salary[data1$gender == 'female']) #41682  
data1$my_prediction2 = ifelse(data1$salary <= maximum_female, 'female', 'male')
```



But this is not learning!

STEPWISE SUPERVISED ML

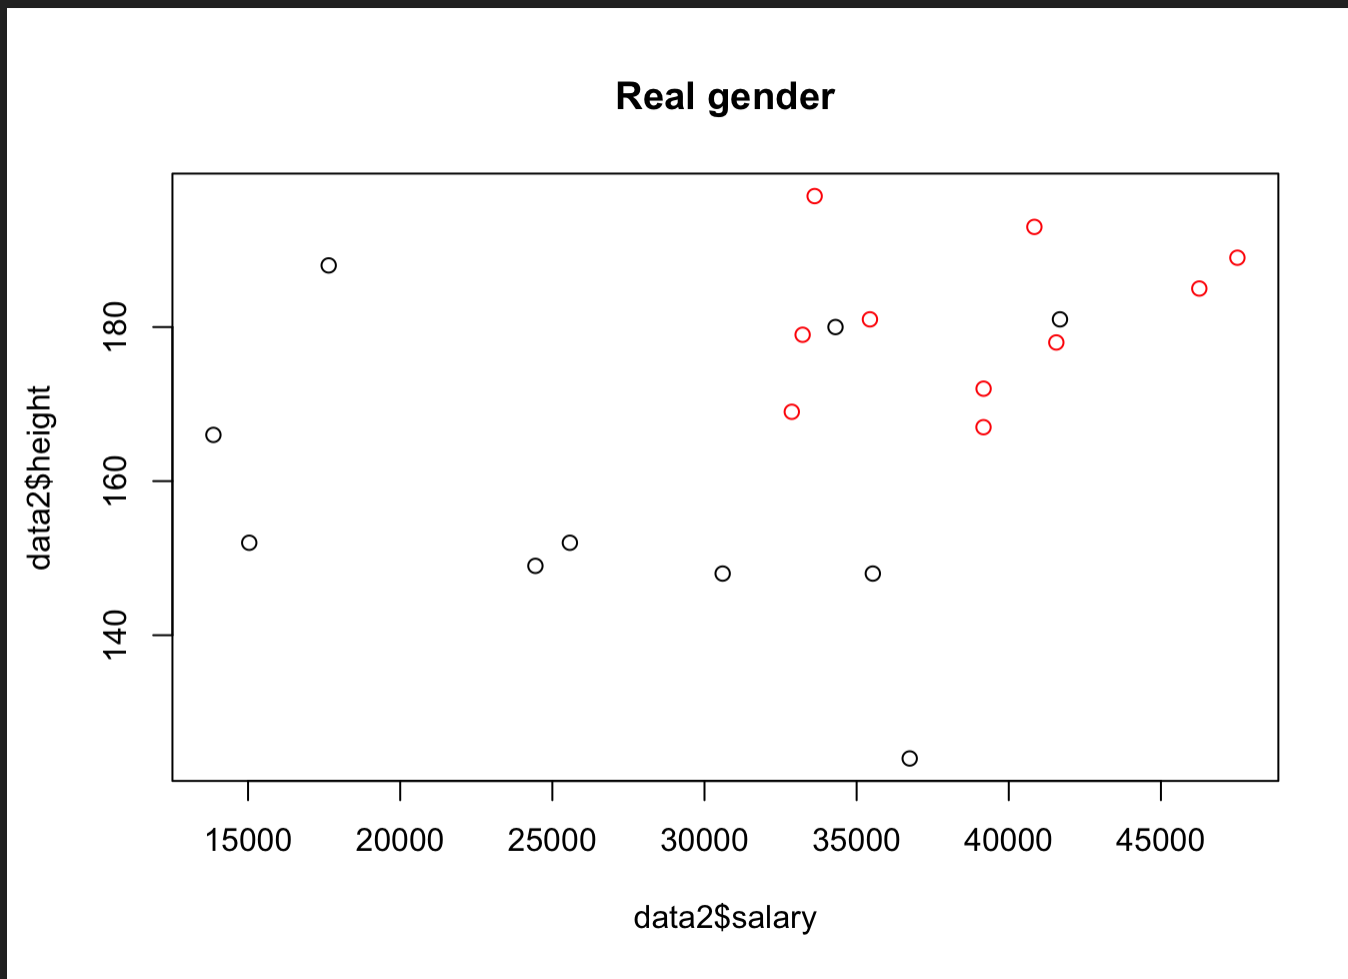
- clarify what `outcome` and `features` are
- determine which classification algorithm to use
- train the model

ENTER: CARET

```
library(caret)
```

- excellent package for ML in R
- well-documented [website](#)
- common interface for [200+ models](#)

CARET IN PRACTICE



CARET IN PRACTICE

```
my_first_model = train(gender ~ .  
                        , data = data2  
                        , method = "svmLinear"  
                        )
```

Now you have trained a model!

you have taught an algorithm to learn to predict gender
from salary & height



But now what?


PUT YOUR MODEL TO USE

Make predictions:

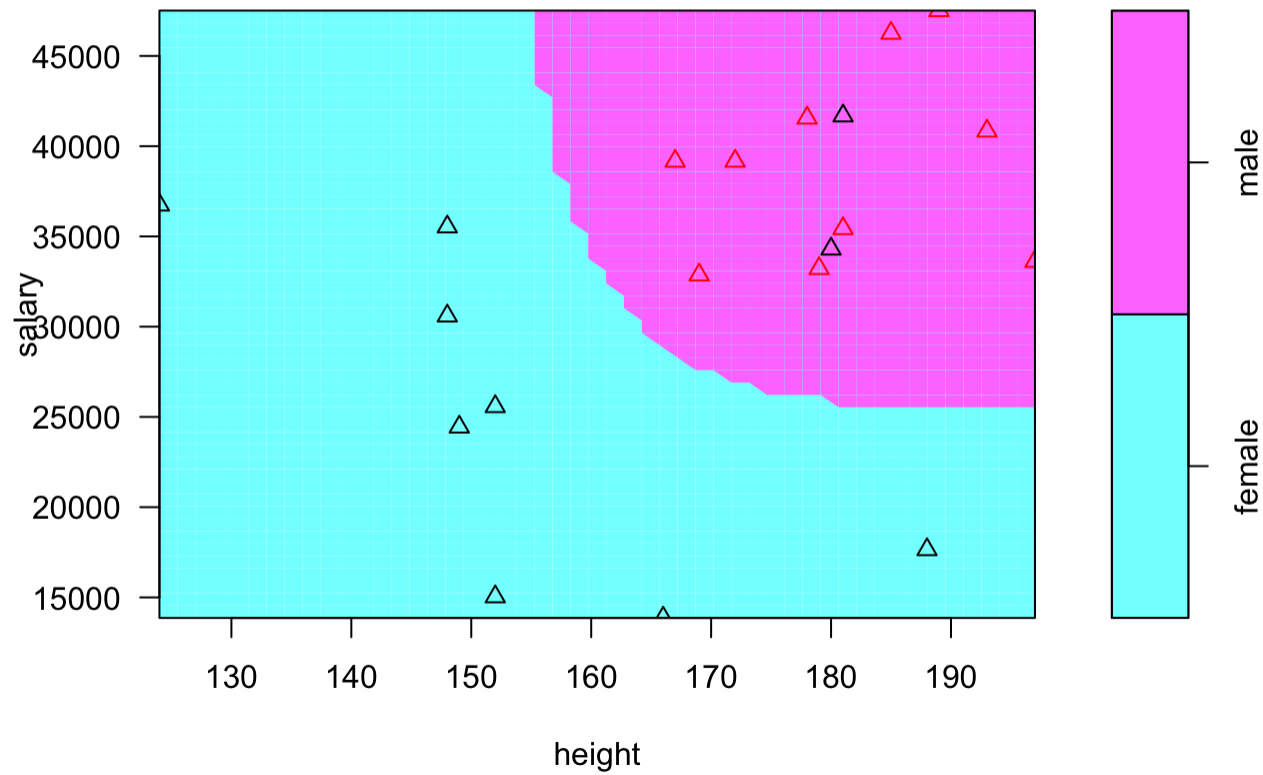
```
data2$model_predictions = predict(my_first_model, data2)
```

	female	male
female	8	2
male	0	10

Algorithm-predicted gender



SVM classification plot



THE KEY CHALLENGE?

Think about what we did...

PROBLEM OF INDUCTIVE BIAS

- remember: we learn from the data
- but what we really want to know is: how does it work on “unseen” data

How to solve this?

KEEP SOME DATA FOR YOURSELF

Train/test split

- split the data (e.g. 80%/20%, 60%/40%)
- use one part as TRAINING SET
- use the other as TEST SET

CARET HELPS!

```
set.seed(1)
in_training = createDataPartition(y = data1$gender
                                   , p = .8
                                   , list = FALSE
                                   )

in_training
```

```
##      Resample1
## [1,]         1
## [2,]         2
## [3,]         4
## [4,]         5
## [5,]         6
## [6,]         7
## [7,]         8
## [8,]        10
## [9,]        12
## [10,]       13
## [11,]       14
## [12,]       15
## [13,]       16
## [14,]       17
## [15,]       18
```

SPLITTING THE DATA

```
training_data = data2[ in_training,]  
test_data = data2[-in_training,]
```

	gender	salary	height
3	male	33225	179
9	male	40841	193
11	female	15039	152
20	female	30597	148

PIPELINE AGAIN

- define outcome (DONE)
- define features (DONE)
- build model (DONE)
 - but this time: on the TRAINING SET
- evaluate model
 - this time: on the TEST SET

Teach the SVM:

```
my_second_model = train(gender ~ .  
                        , data = training_data  
                        , method = "svmLinear"  
                        )
```

Fit/test the SVM:

```
model_predictions = predict(my_second_model, test_data)
```

	female	male
female	2	0
male	0	2

BUT!

- our model might be really dependent on the training data
- we want to be more careful
- Can we do some kind of safeguarding in the training data?

CROSS-VALIDATION

K-fold cross-validation

Iteration 1



Iteration 2



Iteration 3



Iteration 4



Iteration 5



SPECIFYING CV IN CARET

```
training_controls = trainControl(method="cv"  
                                , number = 4  
                                )  
  
my_third_model = train(gender ~ .  
                      , data = training_data  
                      , trControl = training_controls  
                      , method = "svmLinear"  
                      )
```



```
my_third_model
```

```
## Support Vector Machines with Linear Kernel
##
## 16 samples
## 2 predictor
## 2 classes: 'female', 'male'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 12, 12, 12, 12
## Resampling results:
##
##      Accuracy   Kappa
##      0.75      0.5
##
## Tuning parameter 'C' was held constant at a value of 1
```

ASSESS THE CVED MODEL

```
model_predictions = predict(my_third_model, test_data)
```

	female	male
female	2	0
male	0	2

LET'S APPLY THIS!

Fakenews corpus: 1000 fake, 1000 real (data)

	including	ones	information	house	show	security	outcome
1	1	1	1	1	1	1	fake
2	0	0	0	0	0	0	fake
3	0	0	1	2	1	0	fake
1000	0	0	0	0	1	0	fake
1001	1	0	0	0	0	0	real
1002	0	0	0	0	0	0	real
1003	0	0	0	0	0	0	real

PROBLEM

- 1000 fake and 1000 real news items
- only source of information: text
- often fact-checking not available (yet)
- idea: linguistic traces help differentiate fake and real news

```
dim(fake_news_data)
```

```
## [1] 2000 799
```

STEPWISE ML APPROACH

- the outcome variable?
- the features?
- the algorithm?
- the train/test split?
- the training set cross-validation?

MODEL 1

	Model 1
outcome	fake vs real
features	ngram freqs.
algorithm	Linear SVM
train/test	80/20
Cross-val.	10-fold

STEP 1

Partition the data

```
set.seed(2019)
in_training = createDataPartition(y = fake_news_data$outcome
                                   , p = .8 # <-- split value
                                   , list = FALSE
                                   )
training_data = fake_news_data[ in_training,]
test_data = fake_news_data[-in_training,]
```

Training data

Var1	Freq
fake	800
real	800

STEP 2

Define training controls

```
training_controls = trainControl(method="cv"  
                                , number = 10  
                                )
```

STEP 3

Train the model

```
fakenews_model_1 = train(outcome ~ .  
                          , data = training_data  
                          , trControl = training_controls  
                          , method = "svmLinear"  
                          )
```

STEP 4

Fit the model

```
model_1.predictions = predict(fakenews_model_1, test_data)
```

	fake	real
fake	159	41
real	42	158

$$(159+158)/400 = 0.73$$

THE STRENGTH OF CARET...

Let's see whether we can do better

	Model 1	Model 2
outcome	fake vs real	~
features	ngram freqs.	~
algorithm	Linear SVM	~
train/test	80/20	60/40
Cross-val.	10-fold	5-fold

MODEL 2

Step 1: Splitting the data

```
set.seed(2019)
in_training = createDataPartition(y = fake_news_data$outcome
                                   , p = .6 # <-- split value
                                   , list = FALSE
                                   )
training_data = fake_news_data[ in_training,]
test_data = fake_news_data[-in_training,]
```

Step 2: Define training controls

```
training_controls = trainControl(method="cv"  
                                , number = 5  
                                )
```

Step 3: Train the model

```
fakenews_model_2 = train(outcome ~ .  
                          , data = training_data  
                          , trControl = training_controls  
                          , method = "svmLinear"  
                          )
```


Step 4: Fit the model

```
model_2.predictions = predict(fakenews_model_2, test_data)
```

	fake	real
fake	329	71
real	91	309

$$(329+309)/800 = 0.80$$

LOOKING A STEP FURTHER

What's driving the classification?

```
varImp(fakenews_model_1_)
```

```
## ROC curve variable importance
##
##   only 20 most important variables shown (out of 798)
##
##           Importance
## said           100.00
## first           82.70
## last            77.06
## two             73.30
## year            67.09
## years           63.15
## still           58.76
## also           57.69
## three           53.95
## thats          53.91
## one            53.67
```

IMPORTANT FEATURES

“said”

```
tapply(training_data$said, training_data$outcome, mean)
```

```
##          1          0  
## 0.93875 2.97625
```

“first”

```
tapply(training_data$first, training_data$outcome, mean)
```

```
##          1          0  
## 0.48875 1.16000
```

MAKING FULL USE OF CARET

- what if we want to use a different classification algorithm?

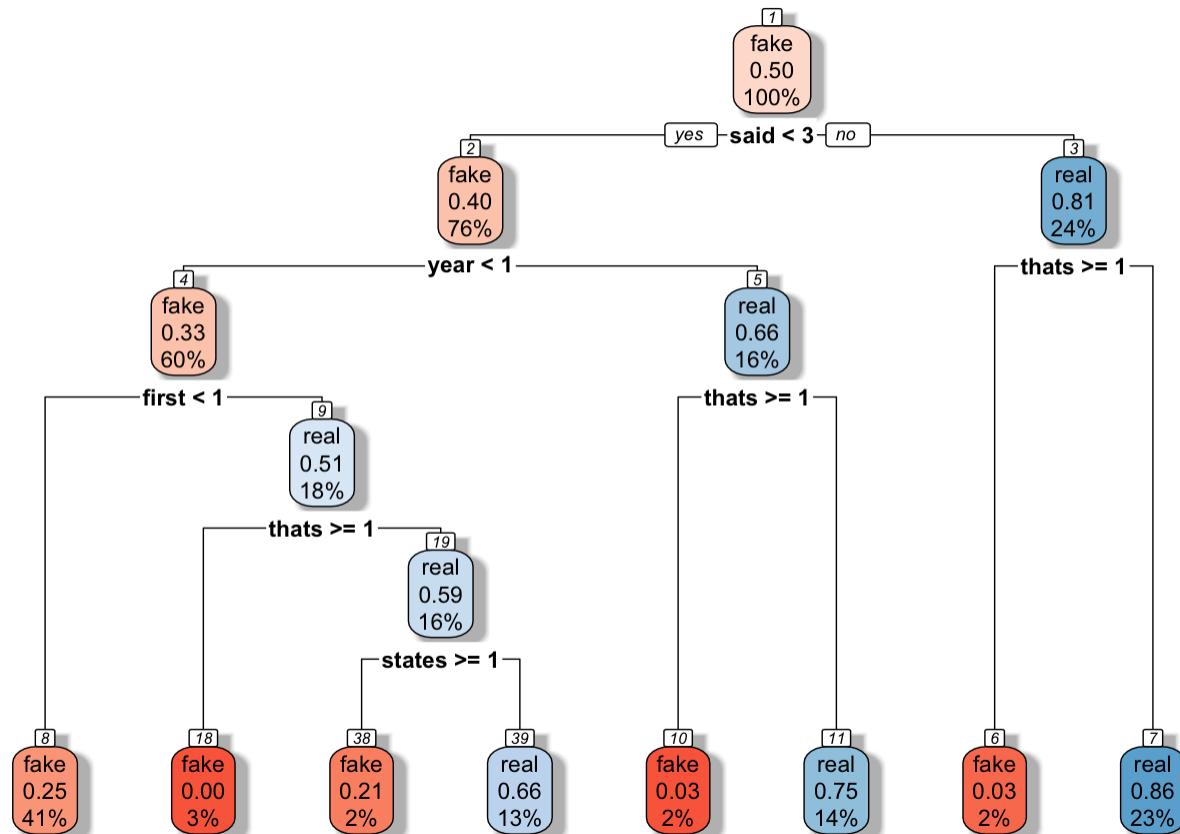
Selection of models ->

<https://topepo.github.io/caret/available-models.html>

INTERMEZZO TO DIFFERENT ALGORITHMS

- Support Vector Machine [video](#)
- Decision Trees
- Random Forests
- worth knowing:
 - Naive Bayes
 - Logistic regression
 - kNN

DECISION TREES



RANDOM FORESTS

- selects random set of training data
 - builds decision tree
- = many trees = forest
- many random trees = random forest
- averaging the trees (voting)

MODEL 3

	Model 1	Model 2	Model 3
outcome	fake vs real	~	~
features	ngram freqs.	~	~
algorithm	Linear SVM	~	Random Forest
train/test	80/20	60/40	70/30
Cross-val.	10-fold	5-fold	2x Repeated 5-fold

MODEL 3

(skipping data splitting here)

```
training_controls = trainControl(method="repeatedcv"  
                                , number = 5  
                                , repeats = 2  
                                )
```

MODEL 3

```
fakenews_model_3 = train(outcome ~ .  
                          , data = training_data  
                          , trControl = training_controls  
                          , method = "ranger"  
                          )
```

MODEL 3

```
## Random Forest
##
## 560 samples
## 798 predictors
## 2 classes: 'fake', 'real'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 2 times)
## Summary of sample sizes: 448, 448, 448, 448, 448, 448, ...
## Resampling results across tuning parameters:
##
##      mtry  splitrule  Accuracy  Kappa
##      2     gini      0.7464286  0.4928571
##      2     extratrees 0.7366071  0.4732143
##     39     gini      0.8250000  0.6500000
##     39     extratrees 0.7901786  0.5803571
```

MODEL 3

Make predictions

```
model_3.predictions = predict(fakenews_model_3, test_data)
```

	fake	real
fake	94	26
real	20	100

$$(90+108)/240 = 0.83$$

RECAP

- Types of machine learning
- Supervised ML
- Cross-validation
- Using caret

OUTLOOK

Tutorial tomorrow

Homework: Replication of fake news classification

Week 7: Machine learning 2

Next week: Unsupervised learning + performance metrics

END