

**MACHINE LEARNING 2**  
**ADVANCED CRIME ANALYSIS**  
**UCL**

**BENNETT KLEINBERG**

**25 FEB 2019**

# MACHINE LEARNING 2

# TODAY

- Recap supervised machine learning
- UNsupervised ML
  - Step-by-step example
- Performance metrics
- Validation and generalisation

# RECAP SUPERVISED ML

- supervised = labeled data
  - classification (e.g. death/alive, fake/real)
  - regression (e.g. income, number of deaths)
- step-wise procedure

# STEPS IN SUPERVISED ML

- clarify what `outcome` and `features` are
- determine which classification algorithm to use
- train the model
  - train/test split
  - cross-validation
- fit the model

# UNSUPERVISED ML

- often we don't have labelled data
- sometimes there are no labels at all
- core idea: finding clusters in the data

```
library(caret)
```


# EXAMPLES

- grouping of online ads
- clusters in crime descriptions
- ...

Practically everywhere.


Clustering reduces your data!





# THE UNSUPERVISED CASE

You know nothing about groups inherent to the data.



# THE K-MEANS IDEA

- separate data in set number of clusters
- find best cluster assignment of observations

# STEPWISE


1. set the number of clusters
2. find best cluster assignment

# 1. NO. OF CLUSTERS

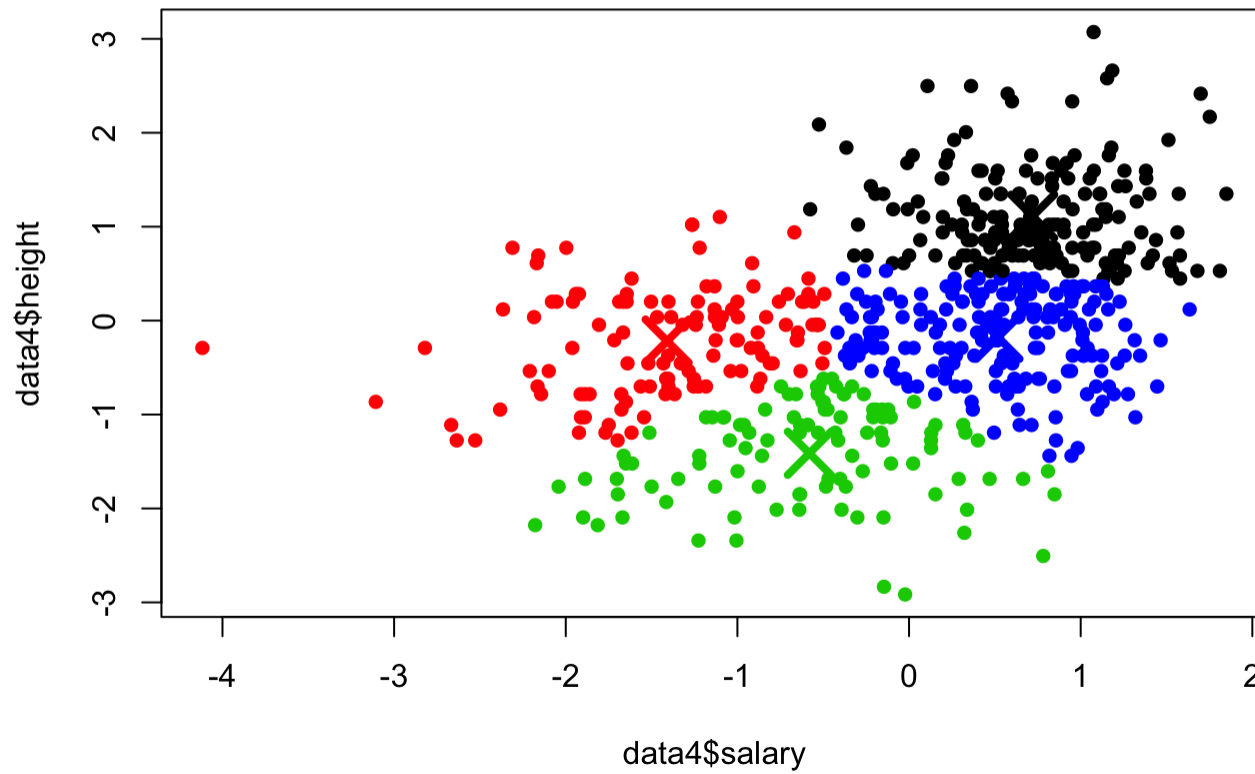
Let's take 4.

```
unsup_model_1 = kmeans(data4  
                        , centers = 4  
                        , nstart = 10  
                        , iter.max = 10)
```

# WHAT'S INSIDE?



Iteration: 2



Cluster plot





# THE K-MEANS ALGORITHM

- find random centers
- assign each observation to its closest center
- optimise for the WSS

**WHAT'S PROBLEMATIC HERE?**

# BUT HOW DO WE KNOW HOW MANY CENTERS?

Possible approach:


- run it for several combinations
- assess the WSS
- determine based on scree-plot

# CLUSTER DETERMINATION

```
wss = numeric()
for(i in 1:20){
  kmeans_model = kmeans(data4, centers = i, iter.max = 20, nstart = 1)
  wss[i] = kmeans_model$tot.withinss
}
```

# SCREE PLOT (ELBOW METHOD)

Look for the inflexion point at center size  $i$ .




# OTHER METHODS TO ESTABLISH $K$


- Silhouette method (cluster fit)
- Gap statistic

See also [this](#) tutorial.

# SILHOUETTE METHOD



# GAP STATISTIC






# CHOOSING $K$

We settle for  $k = 2$

```
unsup_model_final = kmeans(data4  
                             , centers = 2  
                             , nstart = 10  
                             , iter.max = 10)
```

# PLOT THE CLUSTER ASSIGNMENT



# OTHER UNSUPERVISED METHODS

- k-means (today)
- hierarchical clustering
- density clustering

# ISSUES WITH UNSUPERVISED LEARNING


What's lacking?

What can you (not) say?

# CAVEATS OF UNSUP. ML

- there is no “ground truth”
- interpretation/subjectivity
- cluster choice

# INTERPRETATION OF FINDINGS



# INTERPRETATION OF FINDINGS

```
unsup_model_final$centers
```


```
##      salary      height  
## 1 -0.8395549 -0.7457021  
## 2  0.6869085  0.6101199
```

- Cluster 1: low salary, small
- Cluster 2: high salary, tall

Note: we cannot say anything about accuracy.

See the [k-NN model](#).

# INTERPRETATION OF FINDINGS





# INTERPRETATION OF FINDINGS


- subjective
- labelling tricky
- researchers choice!
- be open about this

# CLUSTER CHOICE

What if we chose  $k = 3$ ?

```
km_3 = kmeans(data4, centers = 3, nstart = 10, iter.max = 10)
fviz_cluster(km_3, geom = "point", data = data4)
```

Cluster plot



# CLUSTER CHOICE

What if we chose  $k = 3$ ?

```
km_3$centers
```

```
##      salary      height
## 1  0.7063757  0.8795474
## 2 -1.4058046 -0.5668204
## 3  0.1876933 -0.7256515
```

- Cluster 1: high salary, very tall
- Cluster 2: very low salary, small
- Cluster 3: avg salary, small

# CLUSTER CHOICE

- be open about it
- make all choices transparent
- always share code and data (“least vulnerable” principle)

# PERFORMANCE METRICS FOR CLASSIFICATION TASKS

# FAKE NEWS PROBLEM

# STEP 1: SPLITTING THE DATA

```
set.seed(2019)
in_training = createDataPartition(y = fake_news_data$outcome
                                   , p = .7
                                   , list = FALSE
                                   )
training_data = fake_news_data[ in_training,]
test_data = fake_news_data[-in_training,]
```



## STEP 2: DEFINE TRAINING CONTROLS

```
training_controls = trainControl(method="cv"  
                                , number = 5  
                                , classProbs = T  
                                )
```

## STEP 3: TRAIN THE MODEL

```
fakenews_model = train(outcome ~ .  
                        , data = training_data  
                        , trControl = training_controls  
                        , method = "svmLinear"  
                        )
```

# STEP 4: FIT THE MODEL

```
model.predictions = predict(fakenews_model, test_data)
```

# **YOUR TASK:**

Evaluate the model.

What do you do?

# MODEL EVALUATION

	fake	real
fake	252	48
real	80	220

$$(252+220)/600 = 0.79$$

# **INTERMEZZO**

## **THE CONFUSION MATRIX**

# CONFUSION MATRIX

	Fake	Real
Fake	True positives	False negatives
Real	False positives	True negatives

# CONFUSION MATRIX

- true positives (TP): correctly identified fake ones
- true negatives (TN): correctly identified real ones
- false positives (FP): false accusations
- false negatives (FN): missed fakes



# OKAY: LET'S USE ACCURACIES

$$acc = \frac{(TP+TN)}{N}$$

Any problems with that?

# ACCURACY

Model 1

	<b>Fake</b>	<b>Real</b>
<b>Fake</b>	252	48
<b>Real</b>	80	220

Model 2

	<b>Fake</b>	<b>Real</b>
<b>Fake</b>	290	10
<b>Real</b>	118	182

# PROBLEM WITH ACCURACY

- same accuracy, different confusion matrix
- relies on thresholding idea
- not suitable for comparing models (don't be fooled by the literature!!)

Needed: more nuanced metrics

# BEYOND ACCURACY

```
##           prediction
## reality Fake Real Sum
##   Fake  252   48 300
##   Real   80  220 300
##   Sum   332  268 600
```

```
##           prediction
## reality Fake Real Sum
##   Fake  290   10 300
##   Real  118  182 300
##   Sum  408  192 600
```

# PRECISION

i.e. → how often the prediction is correct when prediction class  $X$

Note: we have two classes, so we get *two* precision values

Formally:

- $Pr_{fake} = \frac{TP}{(TP+FP)}$
- $Pr_{real} = \frac{TN}{(TN+FN)}$

# PRECISION

##	prediction			
##	reality	Fake	Real	Sum
##	Fake	252	48	300
##	Real	80	220	300
##	Sum	332	268	600

- $Pr_{fake} = \frac{252}{332} = 0.76$
- $Pr_{real} = \frac{220}{268} = 0.82$

# COMPARING THE MODELS

	Model 1	Model 2
$acc$	0.79	0.79
$Pr_{fake}$	0.76	0.71
$Pr_{real}$	0.82	0.95

# RECALL

i.e.  $\rightarrow$  how many of class  $X$  is detected

Note: we have two classes, so we get *two* recall values

Also called sensitivity and specificity!

Formally:

- $R_{fake} = \frac{TP}{(TP+FN)}$
- $R_{real} = \frac{TN}{(TN+FP)}$



# RECALL

##		prediction		
##	reality	Fake	Real	Sum
##	Fake	252	48	300
##	Real	80	220	300
##	Sum	332	268	600

- $R_{fake} = \frac{252}{300} = 0.84$
- $R_{real} = \frac{220}{300} = 0.73$

# COMPARING THE MODELS

	Model 1	Model 2
$acc$	0.79	0.79
$Pr_{fake}$	0.76	0.71
$Pr_{real}$	0.82	0.95
$R_{fake}$	0.84	0.97
$R_{real}$	0.73	0.61

# COMBINING PR AND R

The *F1* measure.

Note: we combine Pr and R for each class, so we get *two* F1 measures.

Formally:

- $F1_{fake} = 2 * \frac{Pr_{fake} * R_{fake}}{Pr_{fake} + R_{fake}}$
- $F1_{real} = 2 * \frac{Pr_{real} * R_{real}}{Pr_{real} + R_{real}}$

# F1 MEASURE

##		prediction		
##	reality	Fake	Real	Sum
##	Fake	252	48	300
##	Real	80	220	300
##	Sum	332	268	600

- $F1_{fake} = 2 * \frac{0.76 * 0.84}{0.76 + 0.84} = 2 * \frac{0.64}{1.60} = 0.80$
- $F1_{real} = 2 * \frac{0.82 * 0.73}{0.82 + 0.73} = 0.78$

# COMPARING THE MODELS

	Model 1	Model 2
$acc$	0.79	0.79
$Pr_{fake}$	0.76	0.71
$Pr_{real}$	0.82	0.95
$R_{fake}$	0.84	0.97
$R_{real}$	0.73	0.61
$F1_{fake}$	0.80	0.82
$F1_{real}$	0.78	0.74

# IN CARET

```
confusionMatrix(model.predictions, as.factor(test_data$outcome))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction fake real
##      fake  252    80
##      real   48   220
##
##              Accuracy : 0.7867
##              95% CI : (0.7517, 0.8188)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.5733
##      McNemar's Test P-Value : 0.006143
##
##              Sensitivity : 0.8400
```


# THERE'S MORE

What's actually behind the model's predictions?

Any ideas?


# CLASS PROBABILITIES

Notice anything?

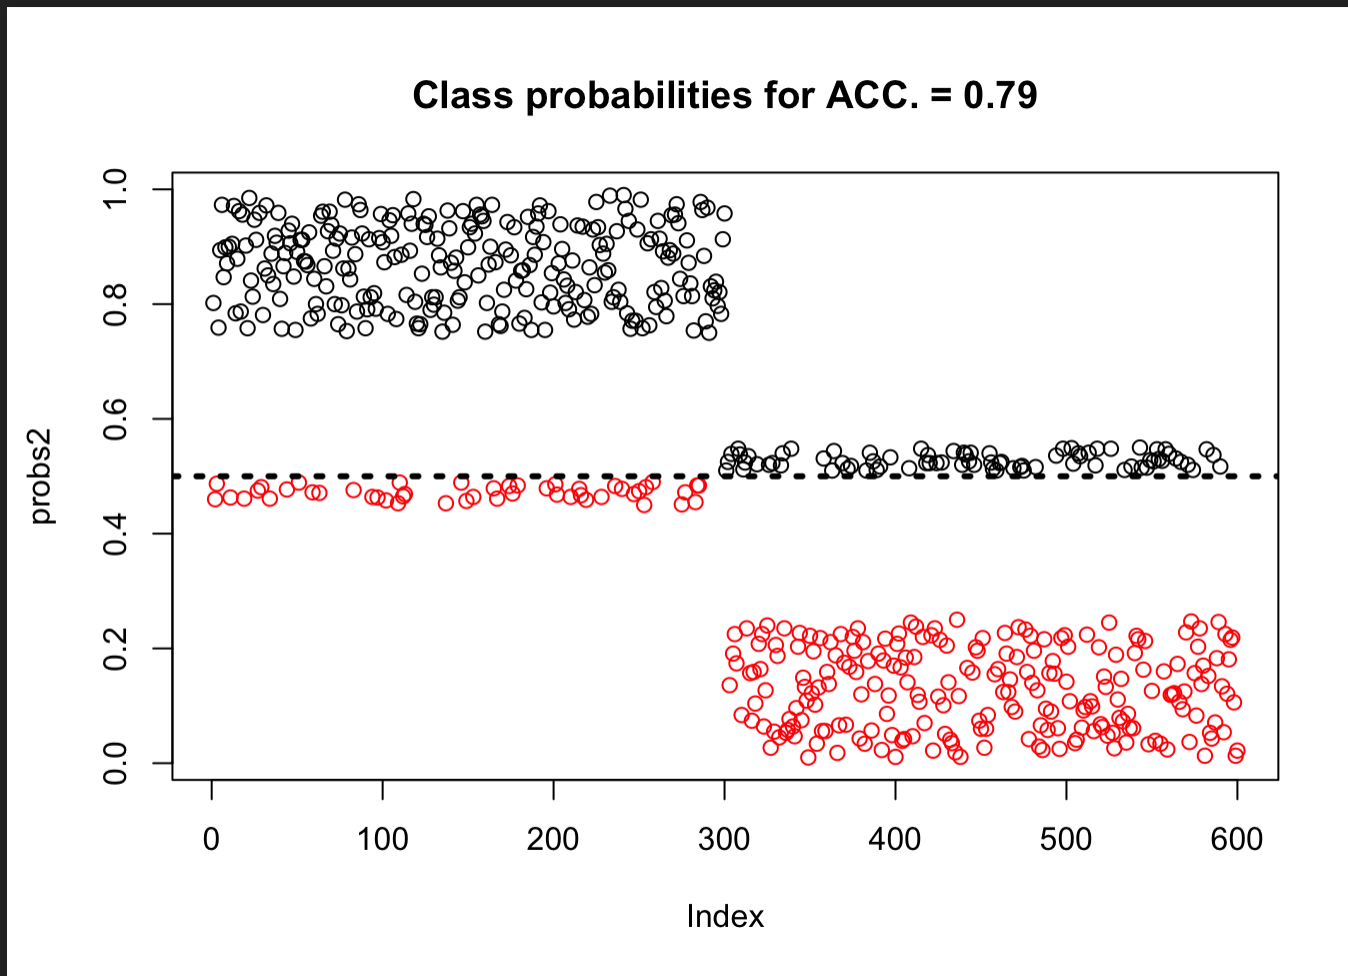




**Class probabilities for ACC. = 0.79**



# THE THRESHOLD PROBLEM



# ISSUE!

- classification threshold little informative
- obscures certainty in judgment

Needed: a representation across all possible values

# THE AREA UNDER THE CURVE (AUC)

Idea:

- plot all observed values (here: class probs)
- y-axis: sensitivity
- x-axis: 1-specificity

# AUC STEP-WISE

```
threshold_1 = probs[1]  
threshold_1
```

```
## [1] 0.6280156
```

```
pred_threshold_1 = ifelse(probs >= threshold_1, 'fake', 'real')  
knitr::kable(table(test_data$outcome, pred_threshold_1))
```

	fake	real
fake	221	79
real	52	248

# SENSITIVITY AND 1-SPECIFICITY

	fake	real
fake	221	79
real	52	248

$$Sens. = 221/300 = 0.74$$

$$Spec. = 248/300 = 0.83$$


$$Sens. = 221/300 = 0.74$$

$$Spec. = 248/300 = 0.83$$

Threshold	Sens.	1-Spec
0.63	0.74	0.17

Do this for every threshold observed.

# .. AND PLOT THE RESULTS:







# QUANTIFY THIS PLOT

```
auc1 = roc(response = test_data$outcome  
           , predictor = probs  
           , ci=T)
```


# WHAT IF WE COMPARE OUR TWO MODELS?




```
plot.roc(auc1, xlim=c(1, 0), legacy.axes = T)
```



**Class probabilities for ACC. = 0.79**



```
auc2 = roc(response = test_data$outcome  
            , predictor = probs2  
            , ci=T)  
  
plot.roc(auc2, xlim=c(1, 0), legacy.axes = T)
```



# AUCS NUMERICALLY

```
#model 1  
roc(response = test_data$outcome , predictor = probs, ci=T)
```

```
##  
## Call:  
## roc.default(response = test_data$outcome, predictor = probs, ci=T)  
##  
## Data: probs in 300 controls (test_data$outcome fake) > 300 cases (test_data$outcome fake)  
## Area under the curve: 0.8521  
## 95% CI: 0.8216-0.8827 (DeLong)
```

```
#model 2  
roc(response = test_data$outcome , predictor = probs2, ci=T)
```

```
##  
## Call:  
## roc.default(response = test_data$outcome, predictor = probs2, ci=T)  
##  
## Data: probs2 in 300 controls (test_data$outcome fake) > 300 cases (test_data$outcome fake)  
## Area under the curve: 0.9573  
## 95% CI: 0.9437-0.971 (DeLong)
```

# RECAP

- Unsupervised ML
- Performance metrics
  - confusion matrix
  - AUC
- Validation and generalisation

# OUTLOOK

Tutorial tomorrow

Week 8: Applied predictive modelling + R Notebooks



**END**